

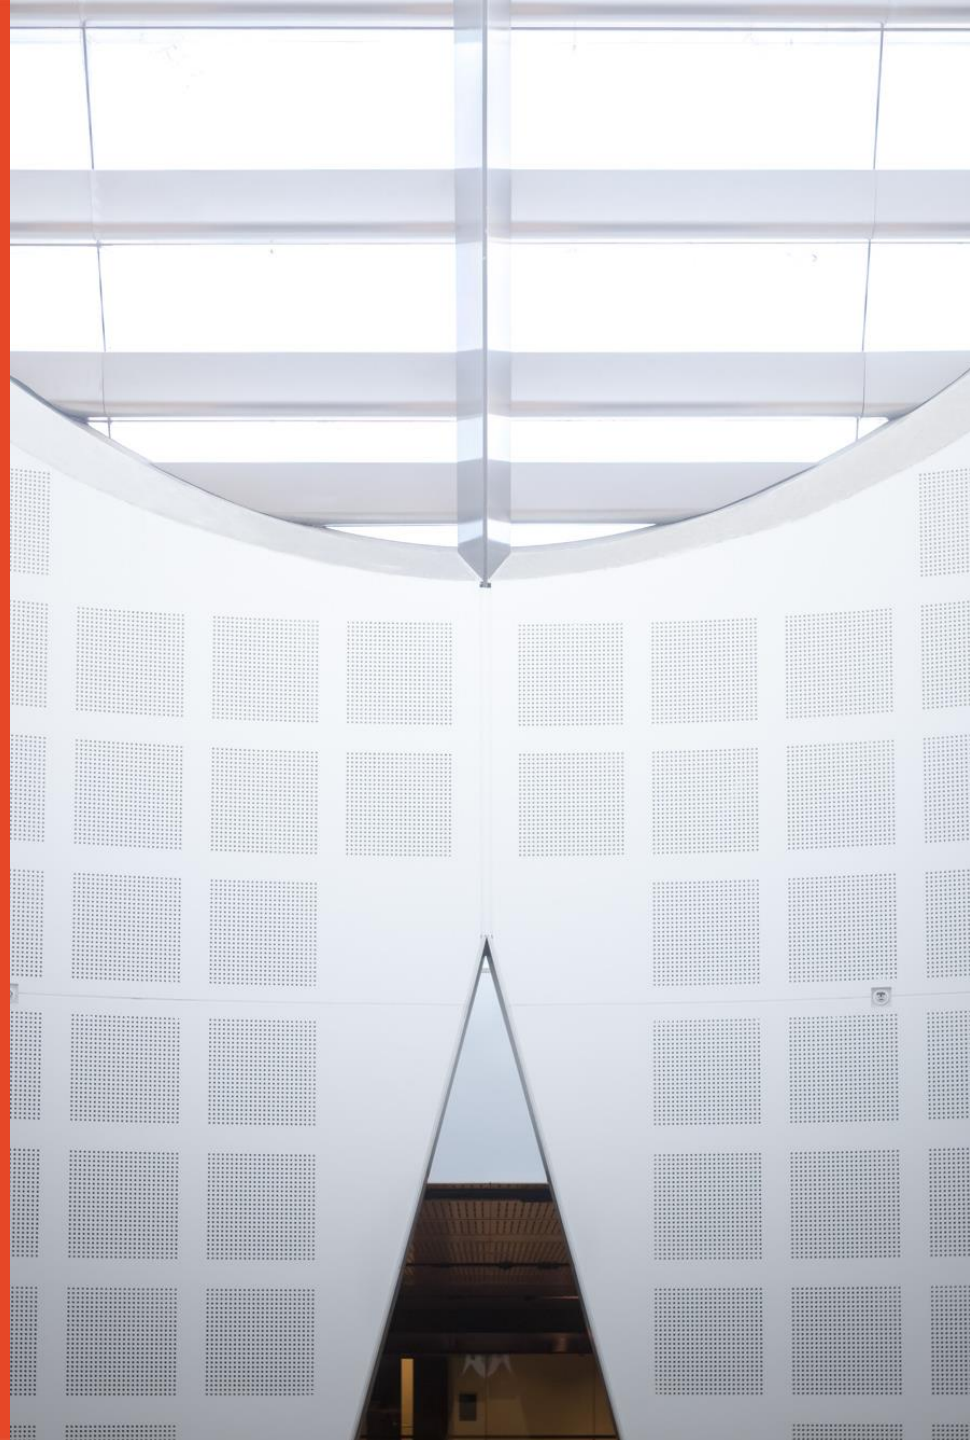
COMP3221: Distributed Systems

Lab 8: Consensus

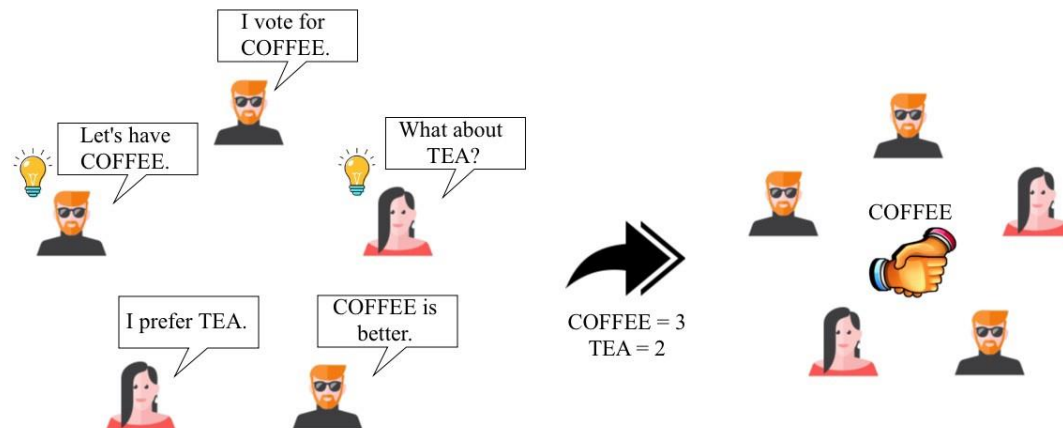
Long Tan Le



THE UNIVERSITY OF
SYDNEY



Consensus in Distributed Systems




- ❑ Consensus is a general agreement on a decision made by the majority of those involved.
- ❑ Three conditions to achieve distributed consensus
 - **Agreement:** All non-faulty processes must agree on the same value.
 - *Every person who didn't leave the conversation (non-faulty process) agrees to drink coffee.*
 - **Termination:** Every process will eventually decide some value.
 - *Everyone in the group promises to make a decision after discussing it for no more than 10 mins.*
 - **Validity:** All the processes must decide on one of the values that were proposed at the beginning.
 - *If they eventually decide on a drink, it'll be one that someone in the group originally proposed.*


Consensus in Blockchain

- ❑ Decentralised Blockchain need a way for users to agree on the current state of the blockchain
- ❑ Several different blockchain consensus mechanisms have been proposed

- Proof of Work:
- Proof of Stake
- Delegate Proof of Stake
- Practical Byzantine Fault Tolerance
- Direct Acyclic Graph
- Proof of Elapsed Time



Force is Power:
Those who own more computing resources govern the network.



Money is Power:
Those who have more money govern the network.



$$R_i = \sum_t \sum_j (R_j * V_{ijt})$$

Reputation is Power:
Those who earn a better reputation and a greater long-term audience base govern the network.

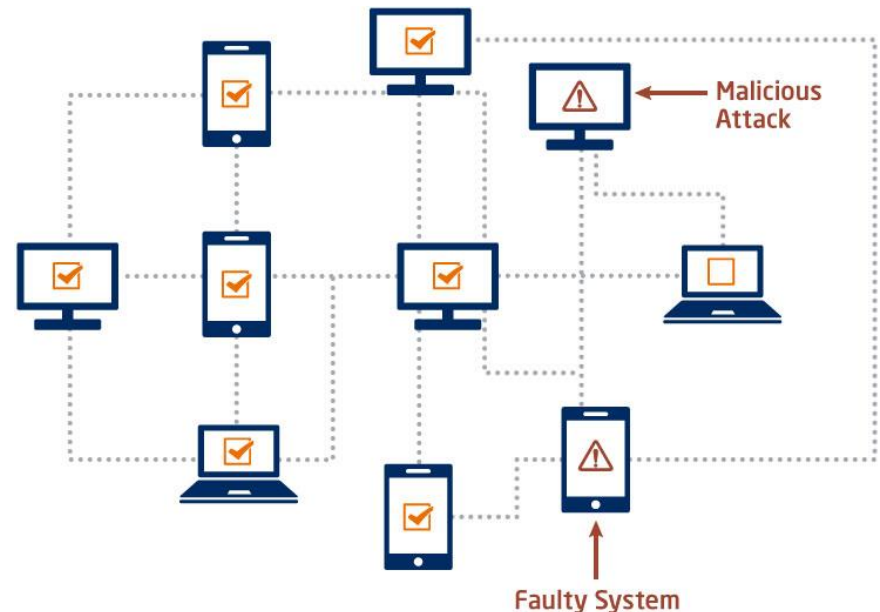
Consensus Problems

❑ Some of nodes might get failed (**Crash fault**) or starts behaving abnormally (**Byzantine Fault**)

- There are n processes, m of which may be faulty ($m < n$)
- The task is to make all the Non-faulty processes agree on some value(s) even **in the presence of the faulty processes**.

❑ Solutions:

- Consensus Without Any Fault (Failure-Free Consensus)
- **Consensus With at most f Crash Faults (Crash-Tolerant Consensus)**
- Consensus With at most f Byzantine Faults



Consensus With at most f Crash Faults


1. Algorithm tolerates at most f failures, out of n nodes ($f < n$)
2. Each process maintains the set of values V proposed by other processes
3. In every round a process: Sends to all other processes the values from V that **it has not sent before**
4. **After $f+1$ rounds** each process decides on the minimum value in V

```
 $P_i$ ::  
var  
     $V$ : set of values initially  $\{v_i\}$ ;  
  
for  $k := 1$  to  $f + 1$  do  
    send  $\{v \in V \mid P_i \text{ has not already sent } v\}$  to all;  
    receive  $S_j$  from all processes  $P_j, j \neq i$ ;  
     $V := V \cup S_j$ ;  
endfor;  
  
 $y := \min(V)$ ;
```


- **Agreement:** Non-faulty processes agree on the same value ($\min(V)$)
- **Termination:** By the code, finish in round $f+1$
- **Validity:** All processes decide on one of the values that were proposed (in V)

Example: $f=1$ failures, $f+1 = 2$ rounds needed

Start




Example: $f=1$ failures, $f+1 = 2$ rounds needed



The failed processor doesn't broadcast its value to all processors

Example: $f=1$ failures, $f+1 = 2$ rounds needed


Round 1



Broadcast all new values to every process

Example: $f=1$ failures, $f+1 = 2$ rounds needed


Round 1



Decide on minimum: two 1s, two 0s -> No agreement

Example: $f=1$ failures, $f+1 = 2$ rounds needed

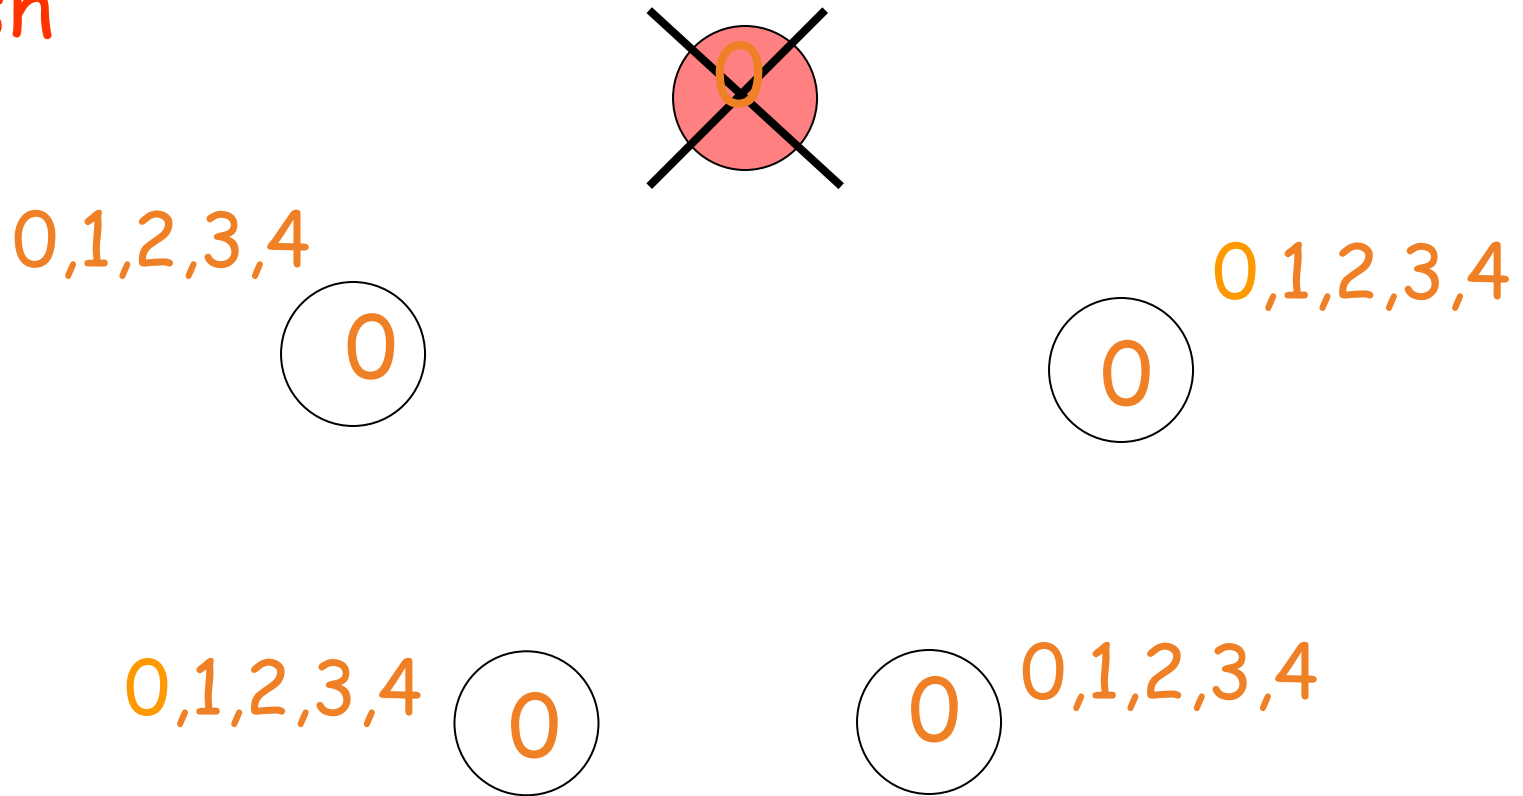
Round 2



Broadcast all new values to everybody

Example: $f=1$ failures, $f+1 = 2$ rounds needed

Finish



Decide on minimum: all 0s -> Agreement